

Programs and Processes

- Program: Executable binary (code and static data)
- Process: A program loaded into memory
 - Program (executable binary with data and text section)
 - Execution state (heap, stack, and processor registers)

Program

```
int foo() {  
    return 0;  
}  
  
int main() {  
    foo();  
    return 0;  
}
```

Process

```
int foo() {  
    return 0;  
}  
  
int main() {  
    foo();  
    return 0;  
}
```

Heap

Stack

Registers

fork()

- A system call that creates a new process identical to the calling one
 - Makes a copy of text, data, stack, and heap
 - Starts executing on that new copy
- Uses of fork()
 - To create a parallel program with multiple processes (E.g. Web server forks a process on each HTTP request)
 - To launch a new program using exec() family of functions (E.g. Linux shell forks an 'ls' process)

fork() example

```
#include <stdio.h>
#include <unistd.h>
int main()
{
    int seq = 0;
    if(fork()==0)
    {
        printf("Child! Seq=%d\n", ++seq);
    }
    else
    {
        printf("Parent! Seq=%d\n", ++seq);
    }
    printf("Both! Seq=%d\n", ++seq);
    return 0;
}
```

```
>> ./a.out
Parent! Seq=1
Both! Seq=2
Child! Seq=1
Both! Seq=2
```


fork() example

```
#include <stdio.h>
#include <unistd.h>
int main()
{
    int seq = 0;
    if(fork()==0)
    {
        printf("Child! Seq=%d\n", ++seq);
    }
    else
    {
        printf("Parent! Seq=%d\n", ++seq);
    }
    printf("Both! Seq=%d\n", ++seq);
    return 0;
}
```

```
>> ./a.out
Parent! Seq=1
Both! Seq=2
Child! Seq=1
Both! Seq=2
```

- Differentiate child and parent using return value of fork()
 - “Child” process return value is 0
 - “Parent” process gets child’s process id number
- **Copies** execution state (not shares)
 - Child copies stack variable seq onto its own stack
 - Child / parent has own copy of seq

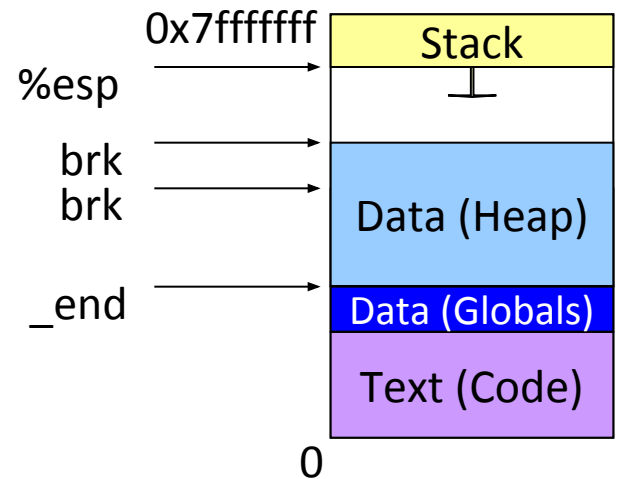
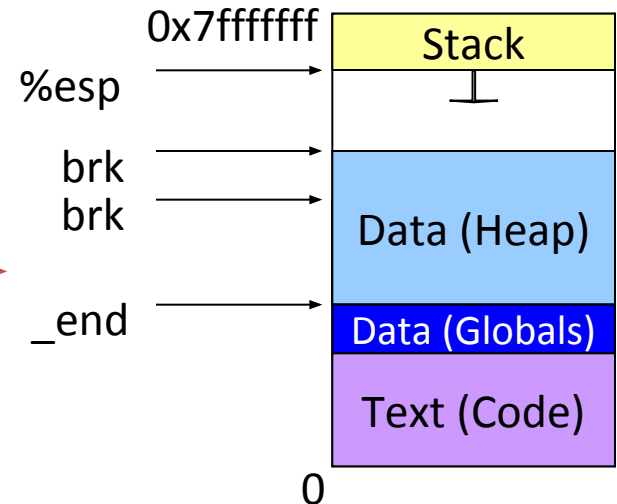
Before fork()



Parent
(original state)

Child
(copy of state)

After fork()



Copy-On-Write

- Inefficient to physically copy memory from parent to child
 - Code (text section) remains identical after fork
 - Even portions of data section, heap, and stack may remain identical after fork
- Copy-On-Write
 - OS memory management policy to **lazily** copy pages only when they are modified
 - Initially map same physical page to child virtual memory space (but in read mode)
 - Write to child virtual page triggers page protection violation (exception)
 - OS handles exception by making physical copy of page and remapping child virtual page to that page

How do Parent / Child Run in Parallel?

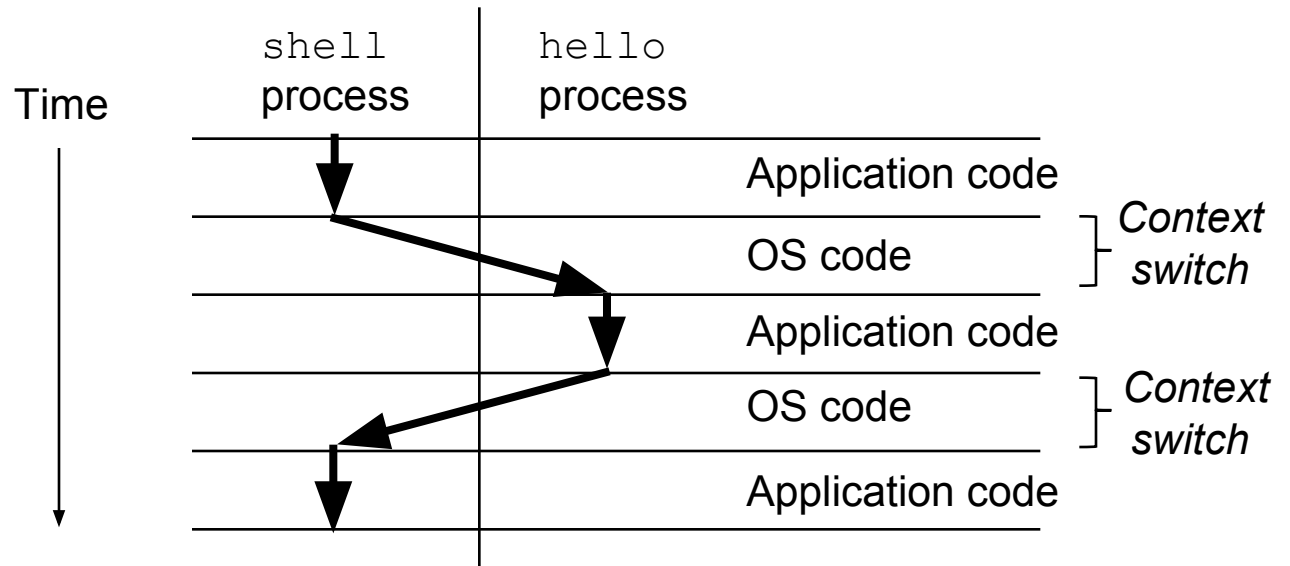
- Obvious answer: by running on different processors on a multiprocessor system
- What if it's a uniprocessor system?
- What if other processors are already busy?
- Answer: by context switching (running each process in short time slots)

Process Context

- Context: set of data that must be saved by a task on interruption and restored on continuation
- Process Context
 - Process registers (including program counter, stack pointer etc.)
 - Memory management info (pointer to page table etc.)
 - I/O status info (pointer to file descriptor table etc.)
 - Process scheduling info (scheduling priority etc.)
 - Process ID, or PID (unique identifier for process)
- Data structure in OS containing process context information is called **process control block (PCB)**

Context Switching

- OS provides the illusion of a dedicated machine per process
- Context switch
 - Saving context of one process, restoring that of another one
 - Processor alternates between executing different processes



Dispatch Mechanism

- OS maintains list of all processes (PCBs)
- Each process has a mode
 - **Running**: Executing on the CPU
 - **Ready**: Waiting to execute on CPU
 - **Blocked**: Waiting for I/O or synchronization with another thread
- Dispatch Loop

```
while (1) {  
    run process for a while;  
    stop process and save its state;  
    load state of another process;  
}
```

How does dispatcher gain control?

- Must change from **user** to **system** mode
 - Problem: Only one CPU, and CPU can only do one thing at a time
 - A user process running means the dispatcher isn't
- **Two ways OS gains control**
- **Exceptions: Events caused by process execution**
 - System calls, page faults, segfault, etc
- **Hardware interrupts: Events external to process**
 - Typing at keyboard, network packet arrivals
 - Control switch to OS via Interrupt Service Routine (ISR)
- **How does OS guarantee it will regain control?**
 - By setting a timer interrupt (allows fair scheduling of processes)

Spawning a New Program

- Combination of `fork()` and `exec(...)`
 - `fork()`: Clone current process
 - `exec(...)`: copy new program on top of current process
- `Exec(...)` family of functions
 - `execl`, `execvp`, `execle`, `execv`, `execve`, `execvp`
 - User space wrappers for the **`execve`** system call
 - Also called the “program loader”
 - Loads in text and data sections of a binary executable
 - Links in any shared objects and perform relocations
 - Sets up stack and starts executing
 - What Linux shell calls when launching a program

execvp() example

```
#include <stdio.h>
#include <unistd.h>
int main()
{
    if(fork()==0)
    {
        char *args[3] = {"ls", "-al", NULL};
        execvp(args[0], args);
        // DOES NOT GET HERE
    }
    else
    {
        printf("Parent!\n");
    }
    printf("Only parent!\n");
    return 0;
}
```

```
>> ./a.out
```

```
Parent!
```

```
Only parent!
```

```
drwx----- 4 wahn UNKNOWN1  4096
Oct 21 08:13 .
```

```
drwxr-xr-x 10 wahn UNKNOWN1  2048
Oct 21 08:13 ..
```

```
-rwxr-xr-x 1 wahn UNKNOWN1  6743
Oct 21 08:12 a.out
```

- execvp never returns since memory is overwritten using another program image

Managing processes (Unix)

- Finding processes
 - 'ps'
- Monitoring Processes
 - 'top'
- Stopping processes
 - 'kill <pid>' (for a soft kill using SIGTERM)
 - 'kill -9 <pid>' (for a hard kill using SIGKILL)
- Procfs (/proc/)

Using 'ps'

- Listing processes associated with this terminal

```
(84) thot $ ps -f
```

UID	PID	PPID	C	STIME	TTY	TIME	CMD
wahn	11848	11847	0	06:27	pts/1	00:00:00	-bash
wahn	15754	11848	0	11:24	pts/1	00:00:00	ps -f

- Listing all processes

```
(85) thot $ ps -ef
```

UID	PID	PPID	C	STIME	TTY	TIME	CMD
root	1	0	0	Aug26	?	00:00:11	/sbin/init
root	2	0	0	Aug26	?	00:00:00	[kthreadd]
root	3	2	0	Aug26	?	00:00:03	[migration/0]
root	4	2	0	Aug26	?	00:00:04	[ksoftirqd/0]
root	5	2	0	Aug26	?	00:00:00	[migration/0]

Using 'kill'

- Killing your own shell

```
(51) thot $ ps -f
```

UID	PID	PPID	C	STIME	TTY	TIME	CMD
wahn	11848	11847	0	06:27	pts/1	00:00:00	-bash
wahn	15904	11848	0	11:36	pts/1	00:00:00	ps -f

```
(52) thot $ kill 11848
```

```
(53) thot $ kill -9 11848
```

```
Connection to thot.cs.pitt.edu closed.
```


Procfs

- File system based export of process information
 - Mounted on /proc/
 - Contains information on every process on the system
 - Organized by PID
 - **/proc/self** exports information for the running process
- Information available
 - Resources in use
 - Resources allowed to use
 - Virtual memory map
 - Cmdline
 - Etc...

Procs usage examples

- Listing open file descriptors for process

```
(84) thot $ ls -l /proc/self/fd
total 0
lrwx----- 1 wahn UNKNOWN1 64 Sep 11 13:37 0 -> /dev/pts/0
lrwx----- 1 wahn UNKNOWN1 64 Sep 11 13:37 1 -> /dev/pts/0
lrwx----- 1 wahn UNKNOWN1 64 Sep 11 13:37 2 -> /dev/pts/0
lr-x----- 1 wahn UNKNOWN1 64 Sep 11 13:37 3 -> /proc/16970/fd
```

- Showing virtual memory map

```
(85) thot $ cat /proc/self/maps
00400000-0040b000 r-xp 00000000 fd:00 2681          /bin/cat
0080a000-0080b000 rw-p 0000a000 fd:00 2681          /bin/cat
0080b000-0082c000 rw-p 00000000 00:00 0            [heap]
34ff600000-34ff78b000 r-xp 00000000 fd:00 131        /lib64/libc-2.12.so
34ff98e000-34ff98f000 rw-p 0018e000 fd:00 131        /lib64/libc-2.12.so
7fffffff000-7fffffff000000 rw-p 00000000 00:00 0    [stack]
```